

COMP9444 Neural Networks and Deep Learning

Quiz 8 Deep RL and Unsupervised Learning

This is an optional quiz to test your understanding of Deep RL and Unsupervised Learning.

1. Write out the steps in the REINFORCE algorithm, making sure to define any symbols you use.

for each trial

run trial and collect states s_t actions a_t and reward r_{total}

for $t = 1$ to length(trial)

$\theta \leftarrow \theta + \eta(r_{\text{total}} - b) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$

end

end

θ = parameters of policy, η = learning rate,

r_{total} = total reward received during trial,

b = baseline (constant), ∇_{θ} = gradient with respect to θ ,

$\pi_{\theta}(a | s)$ = probability of performing action a in state s .

2. In the context of Deep Q-Learning, explain the following:

- a. Experience Replay

The agent(s) choose actions according to their current Q-function, using an ϵ -greedy strategy, and contribute to a central database of experiences in the form (s_t, a_t, r_t, s_{t+1}) . Another thread samples experiences asynchronously from the experience database, and updates the Q-function by gradient descent, to minimize

$$[r_t + \gamma \max_b Q_w(s_{t+1}, b) - Q_w(s_t, a_t)]^2$$

- b. Double Q-Learning

Two sets of Q values are maintained. The current Q-network w is used to select actions, and a slightly older Q-network \bar{w} is used for the target value.

3. What is the Energy function for these architectures:

- a. Boltzmann Machine
- b. Restricted Boltzmann Machine

Remember to define any variables you use.

- a. Boltzmann Machine

$$E(x) = -(\sum_i \sum_j x_i w_{ij} x_j + \sum_i b_i x_i)$$

where x_i = activation of node i (0 or 1)

b. Restricted Boltzmann Machine

$$E(v, h) = -(\sum_i b_i v_i + \sum_j c_j h_j + \sum_{i,j} v_i w_{ij} h_j)$$

where v_i = visible unit activations, h_j = hidden unit activations

4. The Variational Auto-Encoder is trained to maximize

$$\mathbb{E}_{z \sim q_\phi(z | x^{(j)})} [\log p_\theta(x^{(j)} | z)] - D_{\text{KL}}(q_\phi(z | x^{(j)}) || p(z))$$

Briefly state what each of these two terms aims to achieve.

The first term enforces that any sample z drawn from the conditional distribution $q_\phi(z | x^{(j)})$ should, when fed to the decoder, produce something approximating $x^{(j)}$.

The second term encourages the distribution $q_\phi(z | x^{(j)})$ to approximate the Normal distribution $p(z)$ (by minimizing the KL-divergence between the two distributions)

5. Generative Adversarial Networks traditionally made use of a two-player zero-sum game between a Generator G_θ and a Discriminator D_ψ to compute

$$\min_\theta \max_\psi (V(G_\theta, D_\psi))$$

a. Give the formula for $V(G_\theta, D_\psi)$.

$$V(G_\theta, D_\psi) = \mathbb{E}_{x \sim p_{\text{data}}} [\log D_\psi(x)] + \mathbb{E}_{z \sim p_{\text{model}}} [\log(1 - D_\psi(G_\theta(z)))]$$

b. Explain why it may be advantageous to change the GAN algorithm so that the game is no longer zero-sum, and write the formula that the Generator would try to maximize in that case.

The quality of the generated images tends to improve if the Generator instead tries to maximize

$$\mathbb{E}_{z \sim p_{\text{model}}} [\log(D_\psi(G_\theta(z)))]$$

This forces the Generator to put emphasis on improving the poor-quality images, rather than taking the images that are already good and making them slightly better.

6. In the context of GANs, briefly explain what is meant by *mode collapse*, and list three different methods for avoiding it.

Mode collapse is when the Generator produces only a small subset of the desired range of images, or converges to a single image (with minor variations). Methods for avoiding mode collapse include: Conditioning Augmentation, Minibatch Features and Unrolled GANs.